

# The model is the engine. The **harness** is the car.

What makes a good agentic coding harness, and why it matters more than the model.



SECTION 01 / 05

# The Problem

Agents can't remember between sessions, and every obvious fix for it fails.

01

# 2025 was the “Year of the Agent.” It didn't land.

The hype was model-led. In production the model alone fell short, expensively, and the studies agree on the reason: it wasn't the model.

## 95%

of enterprise GenAI pilots showed **no measurable P&L impact**

MIT NANDA, 300+ deployments: the divide “is not driven by model quality” but by approach. The dividing line is memory, workflow integration, and learning.

## 80%

of AI projects fail, **twice** the rate of non-AI software


RAND, from 65 practitioner interviews: of five root causes, only one touches model capability. The rest are problem-alignment, data, and integration.

Across MIT and RAND, one verdict: the missing variable wasn't intelligence, it was the **harness around the model**.

# Models leapfrog every benchmark, then commoditize.

A new frontier model tops every benchmark... until the next does, **days** later. Then it saturates ~81%, open weights reach **parity**, and the benchmark itself gets retired.

## THE TREADMILL • 8 SWE-BENCH RECORDS IN 13 MONTHS, THEN A CEILING



### 🔗 Open weights, now at parity ≈ 80.6% = the frontier

DeepSeek V4-Pro matches the closed SOTA on SWE-bench Verified (Apr '26). Composite gap now ~4 months / 8 ECI pts (Epoch AI, May '26).

### ⚠️ The benchmark itself broke 59% of hard tests flawed

OpenAI **retired SWE-bench Verified** (Feb '26): frontier models had trained on the answers; it's saturated. → SWE-bench Pro.

### 📉 Inference is collapsing ≈ 280× cheaper / 18 mo

GPT-3.5-level tokens: \$20.00 → \$0.07 per M. ~10× every year: "LLMflation."

### "The models are getting commoditized."

Satya Nadella • Microsoft CEO • B62 Pod, Dec 2024

**💎 The model is a commodity, not a moat.**  
Whatever tops the leaderboard this week is matched, undercut and open-weighted within months, and the benchmark itself gets retired.

SECTION 02 / 05


# The Agentic **Harness**

The pilots that stalled tuned the model. The ones that shipped engineered the harness.

02

# Hold the model fixed, change the harness.


Three current frontier models, 2025 to 2026. In each, the model weights never change; only the harness around them does.



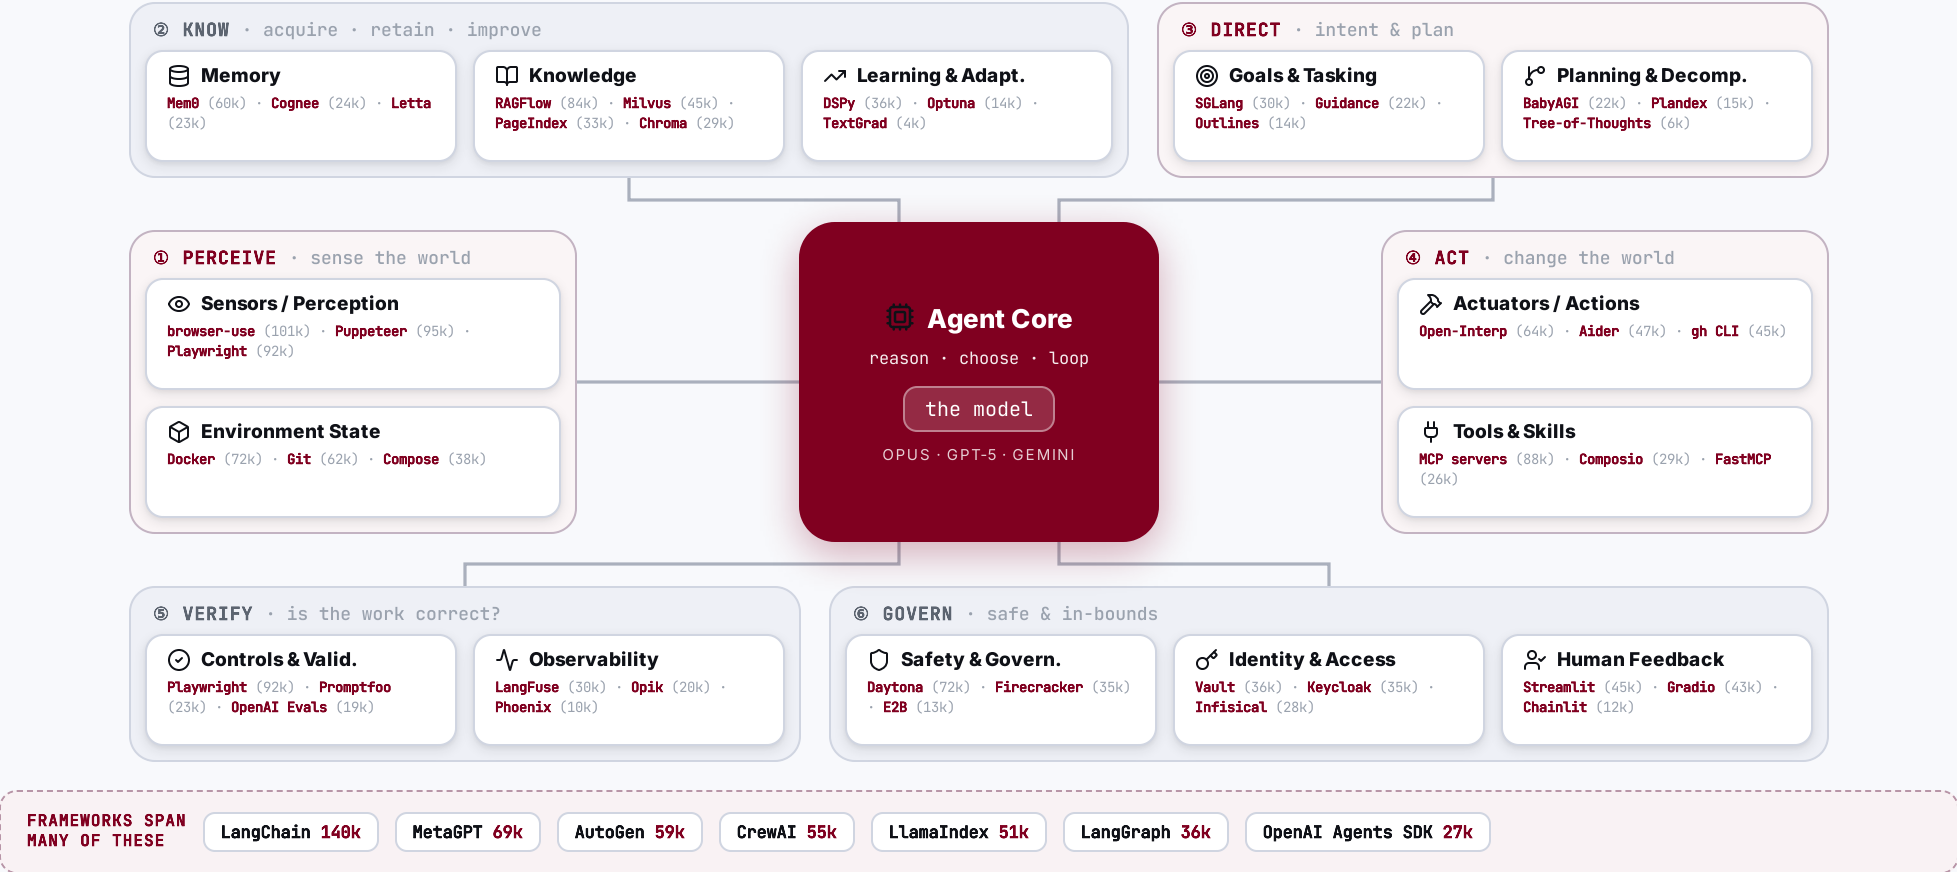
Same weights every time. The harness is the differentiator, **not the model**.

# Performance lives in the tooling around the model, not the model.

The effort runs one way; the leverage runs the other. Teams tune the prompt, but performance comes from the tooling, control, and memory wrapped around the model.



# The agent core, wrapped in 14 capabilities, and who leads each.





# Do you really think software engineering is **dead**?

That map was **50 tools** wrapping a single model: each its own system, all to be wired together and kept alive.

## 50 tools around one model

and only the top few per capability; 80+ ship in the wild



### Each tool is its own system

its own API, config, failure modes and learning curve



### Orchestration is the hard part

wiring 50 moving parts into one reliable, debuggable loop



### Maintenance never stops

versions drift, tools break, the harness needs constant upkeep

## THE LADDER OF ABSTRACTION • every rung was once “the end of programming”

### The agentic harness

2025 • you are here

### Cloud • serverless • DevOps

2010s

### Libraries & frameworks

1990s–2000s

### Compilers & high-level languages

1950s–70s

### Assembly

1950s

### Machine code • punch cards

1940s



## No. We’ve just climbed another rung on the ladder of abstraction.

Software engineering didn’t disappear; building and maintaining this harness is the engineering now.

THE HARDEST LAYER


# Memory

Bolt-on memory is now commodity: Mem0, Letta, vector search. What no tool ships is memory that holds its structure and traceability across refactors, and that is what the harness has to build.



# An agent has four memories. One is volatile, three persist.

The **CoALA** framework (Princeton, TMLR 2024) names them. Everything the agent knows is loaded through the one volatile window to be used.




*One volatile window, three persistent stores. Each fails in its own way, next four slides.*

# No. A bigger context window is not bigger memory.

**Working memory** is the live context window: the only state the model actually sees, and it degrades as you fill it.

## WHERE IT BREAKS



The advertised window is **not** the usable window.

### 59% → 32%

Claude Opus 4.7 loses ~half its multi-needle retrieval from 256K to 1M tokens. Opus 4.6 holds better (93% → 76%).

### Codex too

the GPT-5.x line behind Codex CLI drops ~87% → 74% at long context, a proxy; no per-length figure is published for the Codex-tuned model.

### 11 of 13

models fell below half their accuracy by 32K tokens. Context rot isn't new (NoLiMa, 2025).

→ **The fix:** curate, don't cram. Retrieve only the few high-signal tokens the task needs into a fresh window; treat context as a budget, not a bucket.

# Writing it down isn't memory either.

**Semantic memory** is what the agent knows about your project: the AGENTS.md and CLAUDE.md it reads in at the start of every session.

## WHERE IT BREAKS

OpenAI built a 1M-line product with zero hand-written code. At that scale, they found one big AGENTS.md broke down, four ways:

-  Crowds out the task
-  Too much guidance, none
-  Rots silently
-  Can't be verified

 AGENTS.md

ONE BIG FILE

# Project setup

# Build & test

# Code style

+ 900 more lines • no one can tell which still apply

→ **OpenAI's fix:** they replaced it with a ~100-line table of contents over a structured docs/ tree. Memory has to be a **checkable map**, not a flat file you hope holds.

# Skills tame context bloat, but only if the agent picks the right one.

**Procedural memory** is the agent's reusable know-how: skills, surfaced by progressive disclosure and loaded just-in-time on a matching prompt.

## HOW A SKILL LOADS

AGENT CONTEXT • SYSTEM PROMPT

skill-a

skill-b

skill-c

+ 97

names + descriptions only, always loaded (cheap)

user prompt matches **skill-b** → load it

SKILL-B, JUST-IN-TIME

SKILL.md full instructions

ON MATCH

scripts / references

ON DEMAND

## WHERE IT BREAKS

- ✖

**Discovery rides on one sentence.** Claude picks from 100+ skills using only each description; a vague one never fires.
- 🛒

**The metadata is never free.** Every description permanently shares the context budget.
- 🌀

**It rots silently.** Bundled instructions drift; the only guardrail is "avoid time-sensitive info."
- 📦

**The eager alternative is worse.** Codex **AGENTS.md** and Aider load every turn, no relevance gate.

→ **The fix:** a curated, sharply-described, version-controlled skill set. The procedure layer rewards **structure**, not accumulation.

# The hard part isn't remembering. It's forgetting.

**Episodic memory** is experience from past sessions: the transcripts of what the agent did, and what they taught.

## WHAT IT CAPTURES



### ✓ DISTILLED INTO

#### 📁 KEPT AS EPISODIC MEMORY

💬 Past interactions

# Conversation keywords

✓ What worked

✗ What didn't work

## WHERE IT BREAKS

🕒 **Keeping it all is too costly.** >90% of tokens go to re-stuffing full history each turn.

🔍 **More memory, noisier recall.** The more you retain, the harder the right memory is to find.

📉 **Stale memories mislead.** Old context isn't just bloat; it points the agent the wrong way.

❓ **What to forget is unsolved.** Even scored by importance, recency and relevance, "learned forgetting" is still open (2026).

→ **The fix:** treat retention as a decision, not a default. Choosing what's worth keeping needs a **model of what matters**, which is exactly what no off-the-shelf store has.

# There's a tool for every memory. The recall still breaks.

Map the ecosystem onto CoALA's four memories and every layer is covered. Then the long-term, cross-session benchmarks turn ugly.

## WHAT SHIPS TODAY, PER MEMORY TYPE

**Working memory**  
LLMLingua • context editing • compaction • RAG

**Semantic memory**  
Pinecone • Chroma • Mem0 • Zep / Graphiti • Cognee

**Procedural memory**  
Agent Skills • Cursor Rules • MCP

**Episodic memory**  
Letta • Zep • Mem0 • Windsurf memories

## ⚠ THE MEASURED GAP

~30%

accuracy that commercial assistants and long-context LLMs **lose** recalling across sustained sessions. Up to **60%** on the hard set.

LongMemEval (ICLR 2025) • LoCoMo

*"Current systems oscillate between **hoarding** (store everything, drown in noise) and **amnesia** (compress aggressively, lose rare but vital facts)."*

2026 memory survey, §9

**NOT EVEN MEASURED**

No benchmark tests memory that keeps its **structure and links valid across a refactor**. The field still "lacks a community-standard evaluation harness."



SECTION 03 / 05


# The Metamodel

How an agent armed with skills and rules turns business context into deployable artefacts, every step traceable.

03

# Skills & rules wrap the whole lifecycle, from discovery to operations.


Far beyond DevOps: eight phases from business discovery to operations, each owned by a prefix, all tracing back to one metamodel.



# Discover & Define • make the context real.

WHAT THE AGENTIC ENGINEER DOES HERE

- **Brainstorm & capture**  
ideas, hypotheses
- **Research & scan**  
market, competitors
- **Interview & workshop**  
with the business
- **Model the business**  
vision → objectives
- **Define the domain**  
ubiquitous language



# Architect the solution · turn context into a solution.

WHAT THE AGENTIC ENGINEER DOES HERE

- 


**Decide architecture**  
ADRs, C4, arc42
- 

**Define contracts**  
service, CLI
- 

**Break down function**  
FBS
- 

**Specify & qualify**  
use cases, quality
- 

**Write & review PRDs**  
one per epic



# Execute & Deliver • plan, build, test, ship.

WHAT THE AGENTIC ENGINEER DOES HERE

- 


**Sequence the roadmap**  
epics, phases
- 

**Plan increments**  
atomic steps
- 

**Build & validate**  
tests, checks
- 

**Deploy**  
infra, release
- 

**Operate & fix**  
runbooks, RCA



SECTION 04 / 05


# The Build **Workflow**

From a PRD to merged code: how one engineer drives agents through alignment and verification gates, then lets them run autonomously.

04

# From PRD to merged code: two fresh-session gates, then autonomy.


Business context and discovery are done; the roadmap gives one PRD per epic. The loop below runs per PRD.



Two **fresh-session** checks (grill-me before, peer-review after) keep me and the agent aligned; the loop only runs once the plan is proven autonomous.

# PRD says what must be true; the plan says how, safely.

Two artefacts, one chain: the PRD scopes the epic, the implementation plan breaks it into runnable increments.





# Two fresh agents pressure-test the work, before and after the plan.

Grill-me pulls the thinking out of my head; peer-review checks the document. Both run in a clean session.

## ALIGNMENT · BEFORE THE PLAN

### agent-grill-me

Live Socratic session on the PRD, one question at a time.

- 🔍 Agent asks ~10 focused Q on **implementation & architecture**
- 📖 Vocabulary checked against the **domain glossary**
- 📝 Decisions crystallised into **ADRs** while context is live
- 🧠 Surfaces the **thinking gaps** only questioning reveals

Output: **shared mental model** → the agent then writes the implementation plan.

## VERIFICATION · AFTER THE PLAN

### agent-peer-review

Static, independent scan of the finished plan.

- 🔍 Reads the plan end-to-end, builds a requirement map
- ⚠️ Findings ranked with concrete remediation:
  - critical
  - major
  - normal
  - low
- 🔧 I resolve the ones needing guidance → **amend the plan**

Output: **validated plan** → ready for the autonomy audit.

# The Ralph loop: a fresh agent per increment, until the plan is done.

## WHAT THE LOOP IS

- ↺ A bash script (**ralph.sh**) drives the plan to completion, one increment per iteration.
- ↺ **Fresh agent each iteration** → clean context, no drift or context rot.
- 🚩 Loops while increments are **pending**; stops when every one is **Status: done**.

```
ralph.sh <workspace> \  
  --max-iterations 50 --with-prd --with-push  
# spawns: claude --print < iteration-prompt
```

🔄 ONE ITERATION · PERCEIVE → IMPLEMENT → VALIDATE → COMMIT

Audit must be green before the first run.



### Pick next pending increment

read workspace + plan



### Implement its scope

follow the primary-files list



### Run the test gate

pottery wheel, up to 3 retries



### Commit the increment

advance status → done

↑ repeat with a fresh agent



All increments done → run CI

BUILDING WITH AGENTS • IN PRACTICE

# Dev Tooling

The editor, model, agent-CLI and IDE you build with, open-source or proprietary, and how engineers climb from vibe-coding to real agentic engineering. The harness and the method underneath stay the same.



# Open-source or proprietary: **your setup, your call.**

Three tiers: bring a model, compose a custom harness, or run a full IDE, each available **fully open** or **fully proprietary**. What doesn't change is the method underneath.

|                                                                                                                                     |  <b>OPEN-SOURCE</b> · self-hostable                                                                         |  <b>PROPRIETARY</b> · hosted                                                             |
|-------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|  <b>Full IDE</b><br>ALL-IN-ONE                     | <div><div>VS Code <b>+187k</b></div><div>Zed <b>+86k</b></div><div>Continue <b>+35k</b></div><div>Tabby <b>+34k</b></div></div> <div>Zed standalone · VS Code + ext</div>                    | <div><div>Cursor</div><div>Windsurf</div><div>Antigravity</div><div>Kiro</div></div> <div>closed-source · bundled</div>                                                     |
|  <b>Custom harness</b><br>EDITOR + AGENT           | <div><div>opencode <b>+180k</b></div><div>Cline <b>+64k</b></div><div>Goose <b>+50k</b></div><div>Aider <b>+47k</b></div><div>Crush <b>+26k</b></div></div> <div>open + model-agnostic</div> | <div><div>Claude Code <b>+135k</b></div><div>Codex <b>+94k</b></div><div>Copilot</div><div>Antigravity CLI</div><div>Amp</div></div> <div>proprietary · vendor-hosted</div> |
|  <b>Inference &amp; routing</b><br><b>OPTIONAL</b> | <div><div>Ollama <b>+175k</b></div><div>llama.cpp <b>+118k</b></div><div>vLLM <b>+85k</b></div><div>LiteLLM <b>+52k</b></div></div> <div>self-host · open</div>                              | <div><div>OpenRouter</div><div>Groq</div><div>Hugging Face</div><div>Bedrock</div><div>Vertex AI</div><div>Azure OpenAI</div></div> <div>managed · cloud gateway</div>      |
|  <b>Bring your own model</b><br>THE MODEL          | <div><div>DeepSeek</div><div>Llama</div><div>Qwen</div><div>gpt-oss</div><div>Mistral</div></div> <div>open weights · self-hostable</div>                                                    | <div><div>Claude</div><div>GPT-5</div><div>Gemini 3</div><div>Grok</div></div> <div>frontier · hosted</div>                                                                 |

 **The method**  
THE CONSTANT, EITHER SIDE

spec

 → 

plan

 → 

grill-me

 → 

peer-review


 → 

loop

· artefacts · rules · skills · the metamodel

# From “vibe coding” to *agentic engineering*: five levels.

Autonomy climbs and the harness thickens, but the human moves **up** the stack (prompter → orchestrator), never out of it.



⚠ **Even L5 is \*supervised\* autonomy, not hands-off.** METR: experienced devs were 19% *slower* on mature codebases; SWE-bench Pro tops out ~23%; Ralph’s author: “no way I’d use it on an existing codebase.” The guardrail against Thoughtworks’ “**cognitive debt**” is keeping humans on standards, architecture & final sign-off.

BEFORE YOU GO

# Key takeaways

Ten commandments for the developer, ten for the company. What to actually do on Monday.



# Ten commandments for the **developer**.

The agent is a power tool. The judgement, the architecture and the guardrails are still yours.

**1 Build your own harness, but share what works.**  
No one tells you which IDE or agent to pick; pass the patterns and best practices back to the team.

**2 Treat your agent like a smart junior.**  
Give it tools, steer it, correct it.

**3 Keep a fresh context.**  
Start clean; don't let the window rot.

**4 Agents magnify bad architecture as fast as good.**  
So build the guardrails.

**5 The tooling is hackable.**  
It was built to be bent; adapt it to your workflow.

**6 Start minimal; earn every tool and framework.**  
Hit a wall before you add a tool; run the agent loop bare before you adopt an agentic framework.

**7 Fast to write is not a reason to write.**  
Coding was never the bottleneck; thinking was.

**8 Trust, but verify.**  
Read the diff; every line is yours to defend.

**9 Learn by playing.**  
Build hobby projects, not experiments on company goals.

**10 Refactoring got cheap. Architecture didn't get optional.**  
Cheap to change is not the same as fine to wing it.



Engineer the harness, not just the prompt. The craft moved **up a rung**; it didn't disappear.

# Ten commandments for the **company**.

Building with agents is an organisational capability, not a gadget. Set the conditions.

1

## **Enable the tooling.**

Give every dev agents, sandboxes and access; don't gatekeep the lever.

2

## **Standardize the workflow and artefacts, not the tool.**

Mandate the method; let people pick their own harness.

3

## **AI skips none of your real problems.**

Quality, process, planning, architecture are still there; you just hit the wall faster, and it's bigger.

4

## **Claim AI sovereignty now.**

Don't turn every model and tool into a US dependency.

5

## **Default to open source.**

OSS always ends up the default; don't repeat the Microsoft Office lock-in.

6

## **Govern the data before the agent sees it.**

Decide what company data an agent may ingest; first, not after.

7

## **Budget for maintenance.**

The harness is a living system, not a one-off.

8

## **Pay down cognitive debt.**

Don't let agents quietly erode the team's understanding.

9

## **Measure what must be true, not velocity.**

Track traceability and correctness, not tickets closed.

10

## **The method is the moat, not the model.**

Invest in the harness and the memory; models are commodities.



The model is a commodity. Your edge is the **method and memory** you build around it.



THE TAKEAWAY

# Engineer the **harness**, not just the prompt.

The model is the commodity. Your edge is the method, the memory and the guardrails you build around it.

Same model, better harness

The method is the moat

Standardize the method, swap the parts

Software engineering isn't dead. We just climbed **another rung** on the ladder of abstraction.