


The model is the engine. The **harness** is the car.

What makes a good agentic coding harness — and how clew keeps it coherent.

- 
- The diagram illustrates the agentic coding harness cycle as a series of concentric, rounded rectangular shapes. The innermost shape is a dark blue rounded rectangle. The next layer out is a lighter blue rounded rectangle. The outermost layer is a very light blue rounded rectangle. The cycle is represented by a list of seven steps, each preceded by a small circle. The steps are: SENSE, RETRIEVE, REMEMBER, PLAN, ACT, VALIDATE, and LEARN. The PLAN and ACT steps are highlighted with red circles, while the others have grey circles.
- SENSE
 - RETRIEVE
 - REMEMBER
 - PLAN
 - ACT
 - VALIDATE
 - LEARN


SECTION 01 / 03

The Agentic **Harness**

Why the tooling around the model — not the model — is the performance lever.

01

Hold the model fixed, swap the harness: 42% → 78%.



Claude Opus 4.5 • HAL CORE-Bench • the only thing that changed was the harness

Performance lives in the tooling around the model — not the model.

Ablation studies rank what actually moves agent performance. Model prompt-tuning sits **last**.

1 Tools




2 Middleware



3 Long-term memory



4 System prompt



THE INVERSION


Most teams pour effort into #4 — the prompt. The gains sit in #1-3.

THE MODEL

Table stakes once it is capable. The harness is the differentiator.

One loop, 15 capabilities — a capability system, not a layer cake.

The model fills each step; the harness governs the transitions — and three capabilities govern them all.




GOVERNS
EVERY STEP

Safety & Governance

Identity & Access

Human Feedback

The model is a dot. The harness is the mass around it.



CROSS-CUTTING

Govern every step — they don't belong to one.


 Safety & Governance

 Identity & Access

 Human Feedback


Each wedge is **one loop step** and the capabilities it owns · the three on the right **govern every step**.

Most harnesses only have the shallowest memory — yet memory is the moat.



A bigger context window is not bigger memory.

Retrieval accuracy **rots** as you fill the window — even the 2026 frontier.



The advertised window is **not** the usable window.

59% → 32%

Claude Opus 4.7 loses ~half its multi-needle retrieval from 256K to 1M tokens. Opus 4.6 holds better (93% → 76%).

Codex too

the GPT-5.x line behind Codex CLI drops ~87% → 74% at long context — a proxy; no per-length figure is published for the Codex-tuned model.

11 of 13

models fell below half their accuracy by 32K tokens — context rot isn't new (NoLiMa, 2025).

SECTION 02 / 03

The Metamodel

How an agent armed with skills and rules turns business context into deployable artefacts — every step traceable.

02

From context to code — skills & rules map the whole SDLC.

Every skill carries a prefix; every prefix maps to a phase. The full kit, grouped end to end (utilities aside).

1

Discover & Define

make the context real

- BUSINESS 9
- Vision · Persona · Model canvas · Capability map · Competitive landscape · Value stream · Objective · Process · Quant model
- DISCOVERY 3
- Idea · Research · Workshop
- DOMAIN 3
- Bounded context · Glossary · Domain model

Brainstorm with the business, then model the domain language.

2

Design the Solution

turn context into a solution

- ARCHITECTURE 9
- ADR · Research · Service contract · CLI contract · Structurizr · C4 · PlantUML · UML · arc42
- SPECIFICATION 6
- Functional breakdown · Use case · Quality attributes · PRD · Peer review · Grill-me
- UX 1
- Design system

Specify what to build and decide the architecture — every link traces back to the domain.

3

Execute & Deliver

plan, build, test, ship


- SPECIFICATION 2
- Delivery roadmap · Implementation plan
- QA 5 · pLanned
- Test strategy · Test scenario · Test plan · Acceptance test · Eval harness
- DEVELOPMENT 8
- Agent config · Getting started · Stack guide · Git init · Git commit · Git worktree · PR · Ralph loop
- OPS 3
- Runbook · Bug RCA · Terraform (Exoscale)

Plan increments, validate, then ship and operate — the rest of the SDLC.

Discover & Define — make the context real.

WHAT THE AGENTIC ENGINEER DOES HERE

- **Brainstorm & capture**
ideas, hypotheses
- **Research & scan**
market, competitors
- **Interview & workshop**
with the business
- **Model the business**
vision → objectives
- **Define the domain**
ubiquitous language



Design the Solution — turn context into a solution.

WHAT THE AGENTIC ENGINEER DOES HERE

- 


Decide architecture
ADRs, C4, arc42
- 

Define contracts
service, CLI
- 

Break down function
FBS
- 

Specify & qualify
use cases, quality
- 

Write & review PRDs
one per epic



Execute & Deliver — plan, build, test, ship.

WHAT THE AGENTIC ENGINEER DOES HERE

- 


Sequence the roadmap
epics, phases
- 

Plan increments
atomic steps
- 

Build & validate
tests, checks
- 


Deploy
infra, release
- 

Operate & fix
runbooks, RCA



Everything is related — the metamodel as one graph.

Seven domains, twenty-seven artefacts; solid edges are referential links the engine enforces at write time — dashed are loose post-ship lifecycle.



SECTION 03 / 03

clew

The engine that persists the whole graph: deterministic IDs, write-time integrity, and end-to-end traceability.

03

clew engineers the deepest memory layer — so it stops rotting.

WITHOUT STRUCTURE, MEMORY ROTS



No traceability graph from the docs

No reliable view of how artefacts relate — joins across them are impossible.



No determinism if you trust the LLM

It silently rennumbers IDs; you end up re-verifying everything by hand.



Refactors cost a fortune

Change one ID or name and the effort to update every reference is massive.

CLEW = STRUCTURED LONG-TERM MEMORY



End-to-end traceability

a typed graph generated from the docs themselves



Write-time integrity

deterministic IDs • FK checks • drift detection



Zero token cost per fetch

local file reads • the repo is the source of truth

IDs minted by the app, links checked at write time — never by the LLM.

1 · AGENT ASKS

Call clew via Bash

```
clew new prd --epic E-01
```

The agent never invents an ID — it requests one.

2 · CLEW ENFORCES

Validate & mint

```
E-01 exists? ✓  
FK enforced · → PRD-0001
```

Type-checked against the relationship registry;
collisions rejected.

3 · AGENT WRITES

Reference the ID

```
[PRD-0001](prd-0001.md)
```

Markdown prose cites the stable ID clew returned.



Deterministic — clew mints every ID and checks every link



Guessed — the LLM rennumbers silently and references rot

Ask the graph anything: trace upstream, measure blast radius, prove integrity.

clew trace OBJ-02

OBJ-02
↑ E-01 Epic
↑ C1.2 Capability
↑ P-01 Persona
↑ VISION

Upstream chain — why this exists.

clew impact E-03

E-03 changes →
↓ PRD-0007
↓ QA-SC02
↓ Plan-0007
3 artefacts affected

Downstream blast radius — what breaks.

clew check

scanning references...
bindings ✓
drift ✓
orphans ✓
0 broken references

Integrity audit — deterministic, repeatable.

100% current Every internal reference is up to date, **all the time** — the north-star metric is end-to-end traceability, not work-in-progress.

THE TAKEAWAY

Engineer the **harness**, not just the prompt.

Start with the highest-leverage capability: memory that doesn't rot.

Not a PM tool

Not a wiki

Not a SaaS

The repo is the source of truth

Git for product architecture — manage what must be true, not what's in progress.